

# Data Mining & Machine Learning

## Computer Exercise 5 - Neural Networks & Backpropagation

Steinarr Hrafn Höskuldsson

September 2022

### Section 2.3

The Neural network was trained using 80% of the Iris dataset. The accuracy was 96%, surprisingly good for such a simple model. The confusion matrix was:

$$\text{CM} = \begin{bmatrix} 10 & 0 & 0 \\ 0 & 9 & 0 \\ 0 & 2 & 8 \end{bmatrix}$$

The misclassification rate and the Average error was plotted as a function of iteration number. The plot can be seen in Figure 1. Interestingly, Average error is still trending downwards so it is unclear whether the neural network would improve with more iterations or perhaps overfit.




Figure 1: Average Error and misclassification rate plotted against number of iterations

## Independent

It might be interesting to investigate what would be a good learning rate ( $\eta$ ) for this model and also what might be a good number of iterations( $N$ ). In order to do so a double for loop going through different learning rates and different amounts of iterations was implemented. For each set of learning rate and number of iterations the network was trained and tested on the test set. The resulting accuracies were plotted in pseudocolor plot seen in Figure 2. Interestingly, the best result,

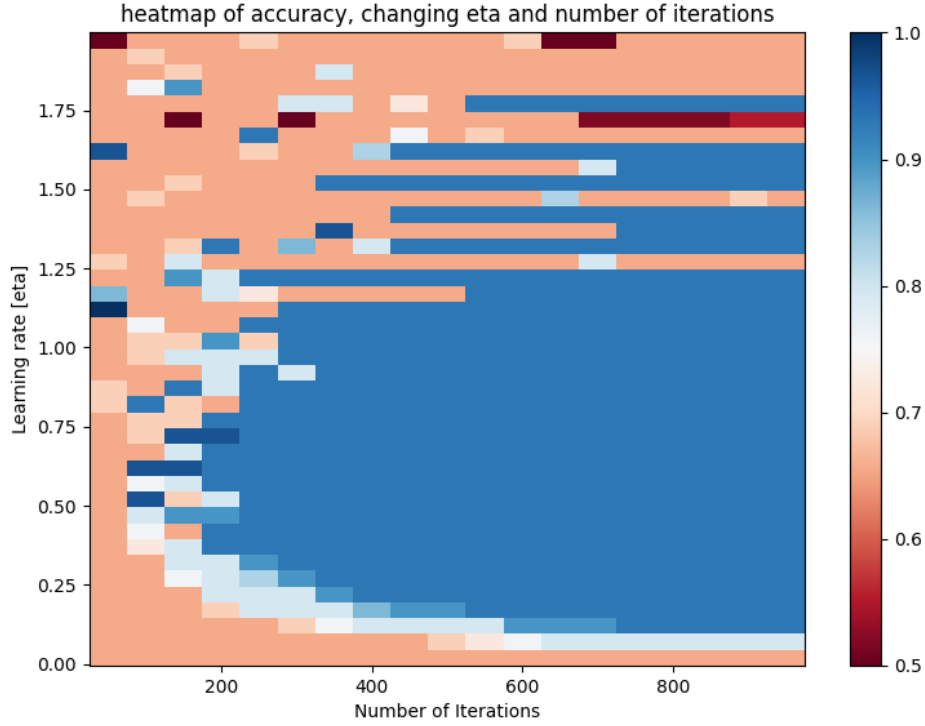


Figure 2: Heatmap of accuracy, changing  $\eta$  and  $N$

100% accuracy, happens when  $\eta = 1.12$  and  $N = 50$ . It is not because those values somehow match a 'training frequency' causing the network to train very effectively. It is because the test set only has 29 samples. And for this particular split of the dataset,  $\eta = 1.12$  and  $N = 50$  just happens to get lucky and correctly guess every test sample correctly.

To confirm this, another for loop was implemented and the above calculation was done for 50 different splits of the dataset into test and train samples. The average of the 50 runs were plotted in a pseudocolor plot seen in Figure 3 and the Standard deviation of the 50 runs can be seen in Figure 4




Figure 3: Heatmap showing average accuracy over 50 runs.

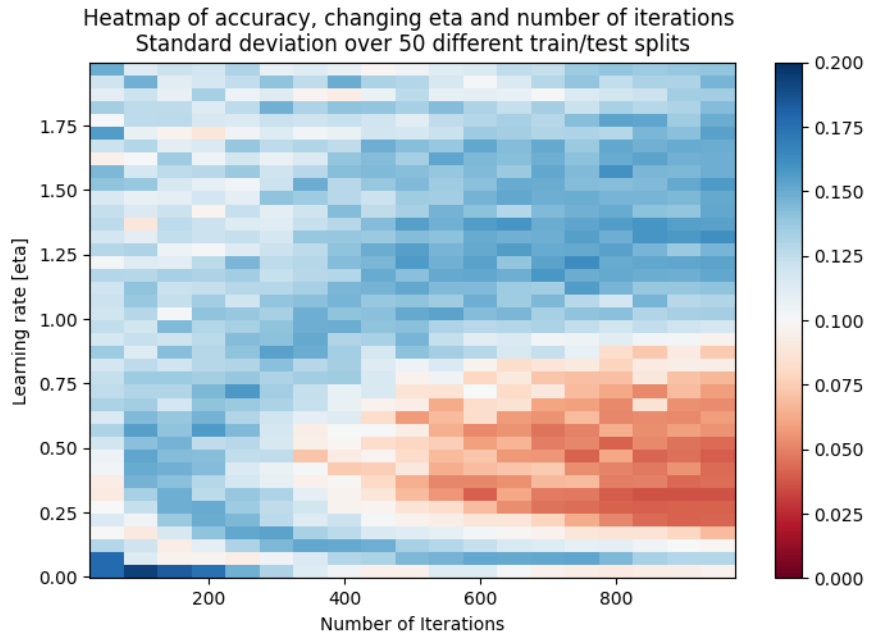


Figure 4: Heatmap showing standard deviation of accuracy over 50 runs.

It is perhaps more interesting to look at a hybrid of the average and standard deviation, A third plot was made showing the average minus twice the standard deviation. This plot can be seen in Figure 5. The highest scoring datapoint in the results was  $\eta = 0.32$  and  $N = 900$  with  $avg - 2 * std = 0.882$  but an interesting point is  $\eta = 0.32$  and  $N = 600$  with  $avg - 2 * std = 0.864$ . If we penalize the number of iterations, for example by dividing by  $\log(N)$  it becomes the highest scoring point. Figure 6 shows  $\log(1000) * (avg - 2 * std) / \log(N)$  plotted for the 50 runs.

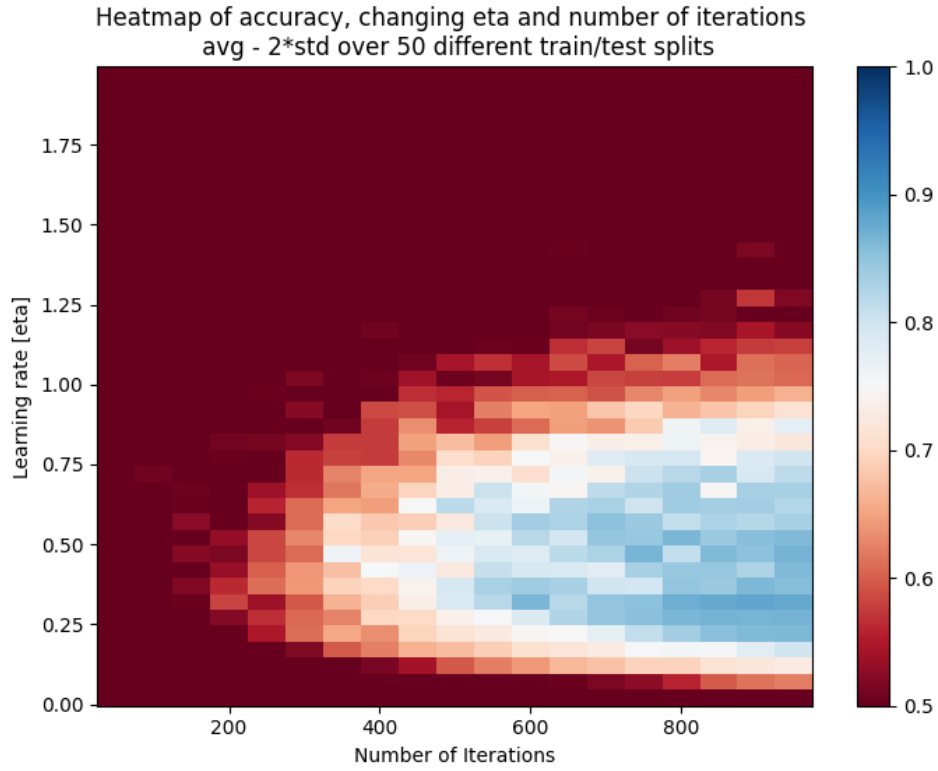


Figure 5: Heatmap showing  $avg - 2 * std$  of accuracy over 50 runs

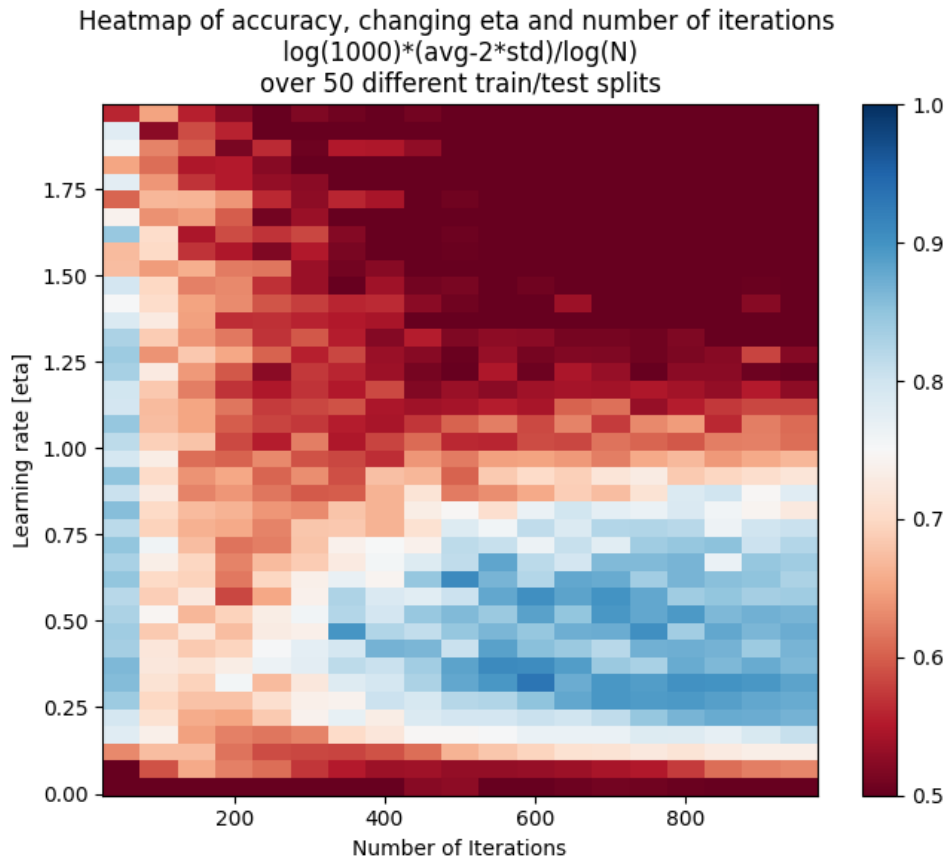


Figure 6: Heatmap showing  $\log(1000) * (\text{avg} - 2 * \text{std}) / \log(N)$  over 50 runs

# Appendix

## A Code

```
1 # author: Steinarr Hrafn
2
3 from typing import Union
4 import numpy as np
5 from matplotlib import pyplot as plt
6 from tools import load_iris, split_train_test
7
8
9 def sigmoid(x: float) -> float:
10     '''
11     Calculate the sigmoid of x
12     '''
13     if isinstance(x, np.ndarray):
14         x[x<-100] = -100
15     elif x < -100:
16         return 0
17     return 1/(1+np.exp(-x))
18
19
20 def d_sigmoid(x: float) -> float:
21     '''
22     Calculate the derivative of the sigmoid of x.
23     '''
24     return sigmoid(x)*(1-sigmoid(x))
25
26
27 def perceptron(
28     x: np.ndarray,
29     w: np.ndarray
30 ) -> Union[float, float]:
31     '''
32     Return the weighted sum of x and w as well as
33     the result of applying the sigmoid activation
34     to the weighted sum
35     '''
36     return np.sum(w*x), sigmoid(np.sum(w*x))
37
38
39 def ffnn(
40     x: np.ndarray,
41     M: int,
42     K: int,
43     W1: np.ndarray,
44     W2: np.ndarray,
45 ) -> Union[np.ndarray, np.ndarray, np.ndarray, np.ndarray, np.ndarray]:
46     '''
47     Computes the output and hidden layer variables for a
48     single hidden layer feed-forward neural network.
49     '''
50     z0 = np.hstack([[1], x])
51     a1 = np.sum(z0*np.transpose(W1), axis=1)
52     z1 = np.hstack([[1], sigmoid(a1)])
53     a2 = np.sum(z1*np.transpose(W2), axis=1)
54     y = sigmoid(a2)
55     # a2 = np.sum(z1*np.transpose(W2[:-1, :])) + W2[-1, :]
```

```

56
57     return y, z0, z1, a1, a2
58
59
60 def backprop(
61     x: np.ndarray,
62     target_y: np.ndarray,
63     M: int,
64     K: int,
65     W1: np.ndarray,
66     W2: np.ndarray
67 ) -> Union[np.ndarray, np.ndarray, np.ndarray]:
68     '''
69     Perform the backpropagation on given weights W1 and W2
70     for the given input pair x, target_y
71     '''
72     y, z0, z1, a1, a2 = ffnn(x, M, K, W1, W2)
73     dk = y - target_y
74     # print( a2.shape, dk.shape, W2.shape, )
75     # print(W2)
76     dj = d_sigmoid(a1)*np.sum(dk*W2[1:, :], axis=1)
77
78     dE1 = dj*z0[..., None]
79     dE2 = dk*z1[..., None]
80     return y, dE1, dE2
81
82
83 def cross_entropy(ts, ys):
84     # assume one hot encoding
85     return -np.sum(ts*np.log(ys) + (1-ts)*np.log(1-ys))
86
87 def one_hot(t, c=3):
88     ret = np.zeros((c))
89     ret[t] = 1
90     return ret
91
92 def hot_one(one):
93     return np.argmax(one, axis=-1)
94
95 def compare_one_hots(y1, y2):
96     return np.argmax(y1) == np.argmax(y2)
97
98 def train_nn(
99     X_train: np.ndarray,
100     t_train: np.ndarray,
101     M: int,
102     K: int,
103     W1: np.ndarray,
104     W2: np.ndarray,
105     iterations: int,
106     eta: float
107 ) -> Union[np.ndarray, np.ndarray, np.ndarray, np.ndarray, np.ndarray]:
108     '''
109     Train a network by:
110     1. forward propagating an input feature through the network
111     2. Calculate the error between the prediction the network
112     made and the actual target
113     3. Backpropagating the error through the network to adjust
114     the weights.
115     '''

```

```

116 # print(t_train)
117 W1tr = W1.copy()
118 W2tr = W2.copy()
119 N = X_train.shape[0]
120 misclassification_rate = []
121 E_total = []
122 guesses = []
123 # loop through iterations
124 for iteration in range(iterations):
125     E = 0
126     misclassifications = 0
127     dE1_total = np.zeros(W1tr.shape)
128     dE2_total = np.zeros(W2tr.shape)
129     # loop through training set
130     for x, y_target in zip(X_train, t_train):
131         y_target = one_hot(y_target)
132         y, dE1, dE2 = backprop(x, y_target, 0, 0, W1tr, W2tr)
133         # y = hot_one(y)
134         dE1_total += dE1
135         dE2_total += dE2
136         if iteration == iterations - 1:
137             guesses.append(hot_one(y))
138             # if iteration == 0 or iteration == iterations-1:
139             #     print(y_target, y, compare_one_hots(y_target, y))
140         E += cross_entropy(y_target, y)
141         misclassifications += not compare_one_hots(y_target, y)
142     W1tr -= eta*dE1_total/N
143     W2tr -= eta*dE2_total/N
144     # guesses = np.array(guesses)
145     E_total.append(E/N)
146     misclassification_rate.append(misclassifications/N)
147 return W1tr, W2tr, E_total, misclassification_rate, guesses
148
149
150 def test_nn(
151     X: np.ndarray,
152     M: int,
153     K: int,
154     W1: np.ndarray,
155     W2: np.ndarray
156 ) -> np.ndarray:
157     """
158     Return the predictions made by a network for all features
159     in the test set X.
160     """
161     return np.array([hot_one(ffnn(x, 0, 0, W1, W2)[0]) for x in X])
162
163
164 if __name__ == '__main__':
165
166     print(f"\n\n{'-' * 20}\n\t Section 1.1 \n")
167     print(f"{sigmoid(0.5)=}")
168     print(f"{d_sigmoid(0.2)=}")
169
170     print(f"\n\n{'-' * 20}\n\t Section 1.2\n")
171     print(f"{perceptron(np.array([1.0, 2.3, 1.9]), np.array([0.2, 0.3, 0.1]))=}")
172     print(f"{perceptron(np.array([0.2, 0.4]), np.array([0.1, 0.4]))=}")
173
174     np.random.seed(34545)
175     print(f"\n\n{'-' * 20}\n\t Section 1.3\n")

```



```

176 features, targets, classes = load_iris()
177 (train_features, train_targets), (test_features, test_targets) = \
178     split_train_test(features, targets)
179 # initialize the random generator to get repeatable results
180 np.random.seed(1234)
181
182 # Take one point:
183 x = train_features[0, :]
184 K = 3 # number of classes
185 M = 10
186 D = 4
187 # Initialize two random weight matrices
188 W1 = 2 * np.random.rand(D + 1, M) - 1
189 W2 = 2 * np.random.rand(M + 1, K) - 1
190 y, z0, z1, a1, a2 = ffnn(x, M, K, W1, W2)
191
192 print(f"{y=}\n{z0=}\n{z1=}\n{a1=}\n{a2=}")
193
194
195 print(f"\n\n{'-' * 20}\n\t Section 1.4\n")
196 # initialize random generator to get predictable results
197 np.random.seed(42)
198
199 K = 3 # number of classes
200 M = 6
201 D = train_features.shape[1]
202
203 x = features[0, :]
204
205 # create one-hot target for the feature
206 target_y = np.zeros(K)
207 target_y[targets[0]] = 1.0
208
209 # Initialize two random weight matrices
210 W1 = 2 * np.random.rand(D + 1, M) - 1
211 W2 = 2 * np.random.rand(M + 1, K) - 1
212
213 y, dE1, dE2 = backprop(x, target_y, M, K, W1, W2)
214
215 print(f"{y=}\n{dE1=}\n{dE2=}")
216
217
218 print(f"\n\n{'-' * 20}\n\t Section 2.1\n")
219 # initialize the random seed to get predictable results
220 np.random.seed(1234)
221
222 K = 3 # number of classes
223 M = 6
224 D = train_features.shape[1]
225
226 # Initialize two random weight matrices
227 W1 = 2 * np.random.rand(D + 1, M) - 1
228 W2 = 2 * np.random.rand(M + 1, K) - 1
229 W1tr, W2tr, Etotol, misclassification_rate, last_guesses = train_nn(
230     train_features[:20, :], train_targets[:20], M, K, W1, W2, 500, 0.1)
231
232 print(f"W1tr = \n{W1tr}\n")
233 print(f"W2tr = \n{W2tr}\n")
234 print(f"Etotol = \n{Etotol[:10]}\n...\n{Etotol[-10:]}\n")
235 print(

```

```

236         f"misclassification_rate = \n{misclassification_rate[:10]}\n...\n{
misclassification_rate[-10:]}\n")
237
238     print(f"last_guesses = \n{last_guesses[:10]}\n...\n{last_guesses[-10:]}\n")
239     print(train_targets[:20])
240
241     print(f"\n\n{'-' * 20}\n\t Section 2.2\n")
242     W1tr, W2tr, Ettotal, misclassification_rate, last_guesses = train_nn(
243         train_features, train_targets, M, K, W1, W2, 500, 0.1)
244     guesses = test_nn(test_features, 0, 0, W1tr, W2tr)
245     print(f"{guesses=}")
246
247     print(f"\n\n{'-' * 20}\n\t Section 2.3\n")
248     accuracy = np.count_nonzero(test_targets==guesses)/len(guesses)
249
250     matrix = np.zeros((3, 3), int)
251     for i, a in enumerate(classes):
252         for j, p in enumerate(classes):
253             matrix[j, i] = np.count_nonzero(guesses[np.where(test_targets == a)] == p
)
254     print(f"{accuracy=: .1%}")
255     print(f"Confusion matrix = \n{matrix}")
256
257     def format_matrix(matrix, environment="bmatrix", formatter=str):
258         """Format a matrix using LaTeX syntax"""
259
260         if not isinstance(matrix, np.ndarray):
261             try:
262                 matrix = np.array(matrix)
263             except Exception:
264                 raise TypeError("Could not convert to Numpy array")
265
266         if len(shape := matrix.shape) == 1:
267             matrix = matrix.reshape(1, shape[0])
268         elif len(shape) > 2:
269             raise ValueError("Array must be 2 dimensional")
270
271         body_lines = [" & ".join(map(formatter, row)) for row in matrix]
272
273         body = "\\\\n".join(body_lines)
274         return f"""\begin{{{environment}}}\n
{body}\n\\end{{{environment}}}\n"""
275     {body}
276     \\end{{{environment}}}\n"""
277     print(format_matrix(matrix))
278
279
280     plt.plot(Ettotal, label="E_total")
281     plt.plot(misclassification_rate, label="misclassification_rate")
282     plt.xlabel("Iterations")
283     plt.legend()
284     plt.show()
285
286
287     print(f"\n\n{'-' * 20}\n\t Independent Section\n")
288
289     # run50 = []
290     # for _ in range(50):
291     #
292     #     (train_features, train_targets), (test_features, test_targets) = \
293     #         split_train_test(features, targets)

```

```

294 #
295 #
296 #     def test_performance(eta, iterations, W1, W2):
297 #         W1tr, W2tr, Etotal, misclassification_rate, last_guesses = train_nn(
train_features[:20, :], train_targets[:20], M, K, W1, W2, iterations, eta)
298 #         guesses = test_nn(test_features, 0, 0, W1tr, W2tr)
299 #         return np.count_nonzero(test_targets == guesses) / len(guesses)
300 #
301 etas = np.arange(0.02, 2.02, 0.05)
302 iterations = np.array(range(50, 1000, 50))
303 #     results = []
304 #     for eta in etas:
305 #         print(eta)
306 #         r = []
307 #         for n in iterations:
308 #             r.append(test_performance(eta, n, W1, W2))
309 #         results.append(r)
310 #     results = np.array(results)
311 #     run50.append(results)
312 # with open("indep_data_50.npy", 'wb') as f:
313 #     np.save(f, np.array(run50))
314 with open("05_backprop/indep_data_50.npy", 'rb') as f:
315     results = np.load(f)
316 results = (np.average(results, axis=0) - 2*np.std(results, axis=0))
317
318 # results = np.std(results, axis=0)
319 print(results)
320 print(etas[results.argmax()//results.shape[1]],
321       iterations[results.argmax() % results.shape[1]])
322
323 fig, ax = plt.subplots()
324
325
326 c = ax.pcolormesh(iterations, etas, results, cmap='RdBu', vmin=0.5, vmax=1,
327                   shading="nearest")
328
329 ax.set_title('Heatmap of accuracy, changing eta and number of iterations \n avg -
330 2*std over 50 different train/test splits')
331
332 ax.set_xlabel("Number of Iterations")
333 ax.set_ylabel("Learning rate [eta]")
334
335 # set the limits of the plot to the limits of the data
336 # ax.axis([x.min(), x.max(), y.min(), y.max()])
337 fig.colorbar(c, ax=ax)
338
339 plt.show()

```