

Data Mining & Machine Learning

Computer Exercise 6 - Neural Networks & PyTorch

Steinarr Hrafn Höskuldsson

September 2022

Section 1

I moved the training to my laptops GPU and added significantly more neurons to both the convolutional layers and the first linear layer. The total accuracy of the network increased from 52% as it was in the training example to 65%. The misclassification rates for each for each category was plotted and can be seen in Figure ??

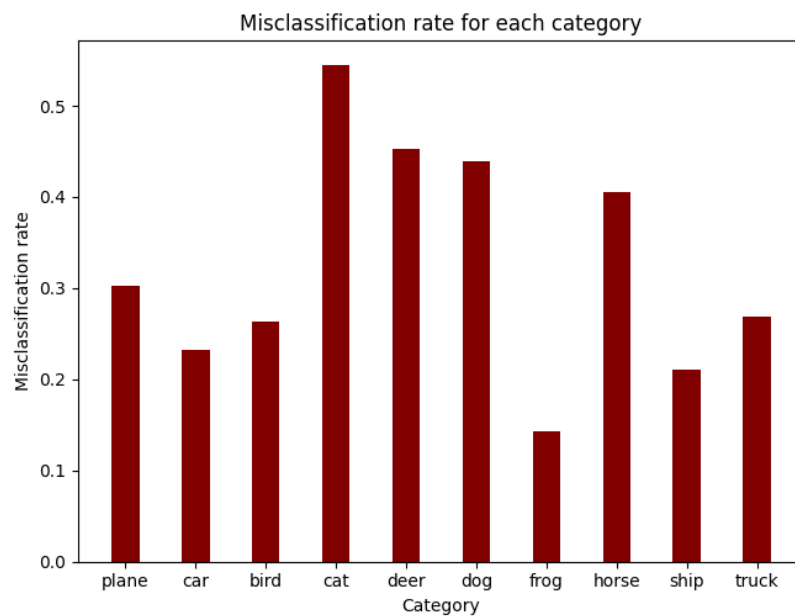


Figure 1: Misclassification rate of each category

The Confusion matrix was calculated.

$$\text{CM} = \begin{bmatrix} & \textit{plane} & \textit{car} & \textit{bird} & \textit{cat} & \textit{deer} & \textit{dog} & \textit{frog} & \textit{horse} & \textit{ship} & \textit{truck} \\ \textit{plane} & 765 & 23 & 16 & 29 & 20 & 3 & 4 & 21 & 74 & 45 \\ \textit{car} & 16 & 839 & 3 & 13 & 4 & 4 & 3 & 4 & 22 & 92 \\ \textit{bird} & 113 & 20 & 350 & 98 & 147 & 110 & 43 & 81 & 19 & 19 \\ \textit{cat} & 33 & 13 & 25 & 499 & 75 & 205 & 21 & 87 & 19 & 23 \\ \textit{deer} & 15 & 4 & 34 & 70 & 595 & 37 & 16 & 206 & 19 & 4 \\ \textit{dog} & 13 & 6 & 17 & 169 & 49 & 607 & 5 & 116 & 5 & 13 \\ \textit{frog} & 7 & 14 & 21 & 156 & 148 & 37 & 578 & 18 & 15 & 6 \\ \textit{horse} & 17 & 0 & 6 & 26 & 41 & 63 & 2 & 817 & 3 & 25 \\ \textit{ship} & 80 & 43 & 2 & 22 & 5 & 10 & 0 & 3 & 787 & 48 \\ \textit{truck} & 38 & 131 & 1 & 14 & 4 & 6 & 2 & 22 & 34 & 748 \end{bmatrix}$$

In the confusion matrix one can spot that the most common misclassifications were mix ups of dog/cat, bird/plane, car/truck, horse/deer/frog which suggests the network might be relying on the background to some extent to classify. Perhaps if the training data also included images with the background extracted the networks could performance could be improved.

The training of the network in PyTorch relied heavily on the AutoGrad which is PyTorch's automatic differentiation engine. It keeps track of operations performed on the dataflow and knows the derivative of those events. When needed, it can apply the chain rule to compute the partial derivative in question.

Section 2

RNN stand for Recursive Neural Network. The recursiveness comes from the fact that the part of the input into the network is the output from the last run. This allows the network to remember what was going on. RNN's are very useful in situations where something is happening over time and the network has to remember what has already happened. A good example of such a situation is working with text. A very important part of having a conversation or completing a sentence is remembering what has previously been said. Other use cases for RNNs are for example video analysis, speech recognition and sequence prediction.

A RNN network was trained based on the example from Pytorch's tutorial. In my opinion the RNN model in Pytorch's example did not perform very well. Although the made up names are usually phonetically viable, they usually did not sound like real names.

Independent Section

All icelandic given names were downloaded from opinskra.is. A short python program was written to parse the valid ones into a .txt file. The Name generation example was then run again and the output generated for Icelandic names was as follows:

- Sari
- Tring
- Ering
- Irin
- Nari
- Arrin
- Rongani
- Rimar

And to a native Icelandic speaker such as myself, none of those sound like an actual name except for the last one, Rimar, which fooled me, I had to look up that it is indeed not a name.

Appendix

A Code

```
1 # author: Steinarr Hrafn
2
3 from typing import Union
4 import numpy as np
5 from matplotlib import pyplot as plt
6 from tools import load_iris, split_train_test
7
8
9 def sigmoid(x: float) -> float:
10     '''
11     Calculate the sigmoid of x
12     '''
13     if isinstance(x, np.ndarray):
14         x[x<-100] = -100
15     elif x < -100:
16         return 0
17     return 1/(1+np.exp(-x))
18
19
20 def d_sigmoid(x: float) -> float:
21     '''
22     Calculate the derivative of the sigmoid of x.
23     '''
24     return sigmoid(x)*(1-sigmoid(x))
25
26
27 def perceptron(
28     x: np.ndarray,
29     w: np.ndarray
30 ) -> Union[float, float]:
31     '''
32     Return the weighted sum of x and w as well as
33     the result of applying the sigmoid activation
34     to the weighted sum
35     '''
36     return np.sum(w*x), sigmoid(np.sum(w*x))
37
38
39 def ffnn(
40     x: np.ndarray,
41     M: int,
42     K: int,
43     W1: np.ndarray,
44     W2: np.ndarray,
45 ) -> Union[np.ndarray, np.ndarray, np.ndarray, np.ndarray, np.ndarray]:
46     '''
47     Computes the output and hidden layer variables for a
48     single hidden layer feed-forward neural network.
49     '''
50     z0 = np.hstack(([1], x))
51     a1 = np.sum(z0*np.transpose(W1), axis=1)
52     z1 = np.hstack(([1], sigmoid(a1)))
53     a2 = np.sum(z1*np.transpose(W2), axis=1)
54     y = sigmoid(a2)
55     # a2 = np.sum(z1*np.transpose(W2[: -1, :])) + W2[-1, :]
```

```

56
57     return y, z0, z1, a1, a2
58
59
60 def backprop(
61     x: np.ndarray,
62     target_y: np.ndarray,
63     M: int,
64     K: int,
65     W1: np.ndarray,
66     W2: np.ndarray
67 ) -> Union[np.ndarray, np.ndarray, np.ndarray]:
68     '''
69     Perform the backpropagation on given weights W1 and W2
70     for the given input pair x, target_y
71     '''
72     y, z0, z1, a1, a2 = ffnn(x, M, K, W1, W2)
73     dk = y - target_y
74     # print( a2.shape, dk.shape, W2.shape, )
75     # print(W2)
76     dj = d_sigmoid(a1)*np.sum(dk*W2[1:, :], axis=1)
77
78     dE1 = dj*z0[..., None]
79     dE2 = dk*z1[..., None]
80     return y, dE1, dE2
81
82
83 def cross_entropy(ts, ys):
84     # assume one hot encoding
85     return -np.sum(ts*np.log(ys) + (1-ts)*np.log(1-ys))
86
87 def one_hot(t, c=3):
88     ret = np.zeros((c))
89     ret[t] = 1
90     return ret
91
92 def hot_one(one):
93     return np.argmax(one, axis=-1)
94
95 def compare_one_hots(y1, y2):
96     return np.argmax(y1) == np.argmax(y2)
97
98 def train_nn(
99     X_train: np.ndarray,
100     t_train: np.ndarray,
101     M: int,
102     K: int,
103     W1: np.ndarray,
104     W2: np.ndarray,
105     iterations: int,
106     eta: float
107 ) -> Union[np.ndarray, np.ndarray, np.ndarray, np.ndarray, np.ndarray]:
108     '''
109     Train a network by:
110     1. forward propagating an input feature through the network
111     2. Calculate the error between the prediction the network
112     made and the actual target
113     3. Backpropagating the error through the network to adjust
114     the weights.
115     '''

```

```

116     # print(t_train)
117     W1tr = W1.copy()
118     W2tr = W2.copy()
119     N = X_train.shape[0]
120     misclassification_rate = []
121     E_total = []
122     guesses = []
123     # loop through iterations
124     for iteration in range(iterations):
125         E = 0
126         misclassifications = 0
127         dE1_total = np.zeros(W1tr.shape)
128         dE2_total = np.zeros(W2tr.shape)
129         # loop through training set
130         for x, y_target in zip(X_train, t_train):
131             y_target = one_hot(y_target)
132             y, dE1, dE2 = backprop(x, y_target, 0, 0, W1tr, W2tr)
133             # y = hot_one(y)
134             dE1_total += dE1
135             dE2_total += dE2
136             if iteration == iterations - 1:
137                 guesses.append(hot_one(y))
138             # if iteration == 0 or iteration == iterations-1:
139                 # print(y_target, y, compare_one_hots(y_target, y))
140             E += cross_entropy(y_target, y)
141             misclassifications += not compare_one_hots(y_target, y)
142         W1tr -= eta*dE1_total/N
143         W2tr -= eta*dE2_total/N
144         # guesses = np.array(guesses)
145         E_total.append(E/N)
146         misclassification_rate.append(misclassifications/N)
147     return W1tr, W2tr, E_total, misclassification_rate, guesses
148
149
150 def test_nn(
151     X: np.ndarray,
152     M: int,
153     K: int,
154     W1: np.ndarray,
155     W2: np.ndarray
156 ) -> np.ndarray:
157     '''
158     Return the predictions made by a network for all features
159     in the test set X.
160     '''
161     return np.array([hot_one(ffnn(x, 0, 0, W1, W2)[0]) for x in X])
162
163
164 if __name__ == '__main__':
165
166     print(f"\n\n{'-' * 20}\n\t Section 1.1 \n")
167     print(f"{sigmoid(0.5)=}")
168     print(f"{d_sigmoid(0.2)=}")
169
170     print(f"\n\n{'-' * 20}\n\t Section 1.2\n")
171     print(f"{perceptron(np.array([1.0, 2.3, 1.9]),np.array([0.2,0.3,0.1]))=}")
172     print(f"{perceptron(np.array([0.2,0.4]),np.array([0.1,0.4]))=}")
173
174     np.random.seed(34545)
175     print(f"\n\n{'-' * 20}\n\t Section 1.3\n")

```

```

176 features, targets, classes = load_iris()
177 (train_features, train_targets), (test_features, test_targets) = \
178     split_train_test(features, targets)
179 # initialize the random generator to get repeatable results
180 np.random.seed(1234)
181
182 # Take one point:
183 x = train_features[0, :]
184 K = 3 # number of classes
185 M = 10
186 D = 4
187 # Initialize two random weight matrices
188 W1 = 2 * np.random.rand(D + 1, M) - 1
189 W2 = 2 * np.random.rand(M + 1, K) - 1
190 y, z0, z1, a1, a2 = ffnn(x, M, K, W1, W2)
191
192 print(f"{y=}\n{z0=}\n{z1=}\n{a1=}\n{a2=}")
193
194
195 print(f"\n\n{'-' * 20}\n\t Section 1.4\n")
196 # initialize random generator to get predictable results
197 np.random.seed(42)
198
199 K = 3 # number of classes
200 M = 6
201 D = train_features.shape[1]
202
203 x = features[0, :]
204
205 # create one-hot target for the feature
206 target_y = np.zeros(K)
207 target_y[targets[0]] = 1.0
208
209 # Initialize two random weight matrices
210 W1 = 2 * np.random.rand(D + 1, M) - 1
211 W2 = 2 * np.random.rand(M + 1, K) - 1
212
213 y, dE1, dE2 = backprop(x, target_y, M, K, W1, W2)
214
215 print(f"{y=}\n{dE1=}\n{dE2=}")
216
217
218 print(f"\n\n{'-' * 20}\n\t Section 2.1\n")
219 # initialize the random seed to get predictable results
220 np.random.seed(1234)
221
222 K = 3 # number of classes
223 M = 6
224 D = train_features.shape[1]
225
226 # Initialize two random weight matrices
227 W1 = 2 * np.random.rand(D + 1, M) - 1
228 W2 = 2 * np.random.rand(M + 1, K) - 1
229 W1tr, W2tr, Etotal, misclassification_rate, last_guesses = train_nn(
230     train_features[:20, :], train_targets[:20], M, K, W1, W2, 500, 0.1)
231
232 print(f"W1tr = \n{W1tr}\n")
233 print(f"W2tr = \n{W2tr}\n")
234 print(f"Etotal = \n{Etotal[:10]}\n...\n{Etotal[-10:]}\n")
235 print(

```

```

236         f"misclassification_rate = \n{misclassification_rate[:10]}\n...\n{
misclassification_rate[-10:]}\n")
237
238     print(f"last_guesses = \n{last_guesses[:10]}\n...\n{last_guesses[-10:]}\n")
239     print(train_targets[:20])
240
241     print(f"\n\n{'-' * 20}\n\t Section 2.2\n")
242     W1tr, W2tr, Ettotal, misclassification_rate, last_guesses = train_nn(
243         train_features, train_targets, M, K, W1, W2, 500, 0.1)
244     guesses = test_nn(test_features, 0, 0, W1tr, W2tr)
245     print(f"{guesses=}")
246
247     print(f"\n\n{'-' * 20}\n\t Section 2.3\n")
248     accuracy = np.count_nonzero(test_targets==guesses)/len(guesses)
249
250     matrix = np.zeros((3, 3), int)
251     for i, a in enumerate(classes):
252         for j, p in enumerate(classes):
253             matrix[j, i] = np.count_nonzero(guesses[np.where(test_targets == a)] == p)
254     print(f"{accuracy=: .1%}")
255     print(f"Confusion matrix = \n{matrix}")
256
257     def format_matrix(matrix, environment="bmatrix", formatter=str):
258         """Format a matrix using LaTeX syntax"""
259
260         if not isinstance(matrix, np.ndarray):
261             try:
262                 matrix = np.array(matrix)
263             except Exception:
264                 raise TypeError("Could not convert to Numpy array")
265
266         if len(shape := matrix.shape) == 1:
267             matrix = matrix.reshape(1, shape[0])
268         elif len(shape) > 2:
269             raise ValueError("Array must be 2 dimensional")
270
271         body_lines = [" & ".join(map(formatter, row)) for row in matrix]
272
273         body = "\\\\\\".join(body_lines)
274         return f"\\begin{{{environment}}}"
275     {body}
276     \\end{{{environment}}}"
277     print(format_matrix(matrix))
278
279
280     plt.plot(Etotal, label="E_total")
281     plt.plot(misclassification_rate, label="misclassification_rate")
282     plt.xlabel("Iterations")
283     plt.legend()
284     plt.show()
285
286
287     print(f"\n\n{'-' * 20}\n\t Independent Section\n")
288
289     # run50 = []
290     # for _ in range(50):
291     #
292     #     (train_features, train_targets), (test_features, test_targets) = \
293     #         split_train_test(features, targets)
294     #

```



```

295 #
296 #     def test_performance(eta, iterations, W1, W2):
297 #         W1tr, W2tr, Etotat, misclassification_rate, last_guesses = train_nn(
train_features[:20, :], train_targets[:20], M, K, W1, W2, iterations, eta)
298 #         guesses = test_nn(test_features, 0, 0, W1tr, W2tr)
299 #         return np.count_nonzero(test_targets == guesses) / len(guesses)
300 #
301 etas = np.arange(0.02, 2.02, 0.05)
302 iterations = np.array(range(50, 1000, 50))
303 #     results = []
304 #     for eta in etas:
305 #         print(eta)
306 #         r = []
307 #         for n in iterations:
308 #             r.append(test_performance(eta, n, W1, W2))
309 #         results.append(r)
310 #     results = np.array(results)
311 #     run50.append(results)
312 # with open("indep_data_50.npy", 'wb') as f:
313 #     np.save(f, np.array(run50))
314 with open("05_backprop/indep_data_50.npy", 'rb') as f:
315     results = np.load(f)
316 results = (np.average(results, axis=0) - 2*np.std(results, axis=0))
317
318 # results = np.std(results, axis=0)
319 print(results)
320 print(etas[results.argmax()//results.shape[1]],
321       iterations[results.argmax() % results.shape[1]])
322
323 fig, ax = plt.subplots()
324
325
326 c = ax.pcolormesh(iterations, etas, results, cmap='RdBu', vmin=0.5, vmax=1, shading="
nearest")
327
328 ax.set_title('Heatmap of accuracy, changing eta and number of iterations \n avg - 2*std
over 50 different train/test splits')
329
330 ax.set_xlabel("Number of Iterations")
331 ax.set_ylabel("Learning rate [eta]")
332
333 # set the limits of the plot to the limits of the data
334 # ax.axis([x.min(), x.max(), y.min(), y.max()])
335 fig.colorbar(c, ax=ax)
336
337 plt.show()

```